

classification_Xinming

Xinming Liu

11/27/2020

Set A - Debate 1

```
## SetA Preprocessing
users <- read.csv("user_setA/users.csv", stringsAsFactors = FALSE)
#head(a.users)
tweets <- read.csv("user_setA/tweets_debate3.csv", stringsAsFactors = FALSE)
# tweets2 <- read.csv("user_setA/tweets_debate2.csv", stringsAsFactors = FALSE)
# tweets3 <- read.csv("user_setA/tweets_debate3.csv", stringsAsFactors = FALSE)
# tweets4 <- read.csv("user_setA/tweets_debateVP.csv", stringsAsFactors = FALSE)
# tweets <- rbind(tweets1, tweets2, tweets3, tweets4)
#head(tweets)
tweets<-tweets%>%
  mutate(userID=as.numeric((userID)))
data <- data.table(users, key="userID")[
  data.table(tweets, key="userID"),
  allow.cartesian=TRUE
]
data <- subset(data, party=="D" | party=="R")
#data <- subset(data, state_code=="PA")
#head(data)
use <- data.frame(data$text, factor(data$party), stringsAsFactors = FALSE)
colnames(use) <- c("text", "party")
#head(use)
```

```
## Create DocumentTermMatrix
corpus <- Corpus(VectorSource(use$text))
corpus = clean_corpus(corpus)
#td.mat = TermDocumentMatrix(corpus)
dt.mat = DocumentTermMatrix(corpus)
## dt.mat is not a matrix here
```

```
## Feature words extraction (this may create NA values) due to limited memory
dt.mat.use = removeSparseTerms(dt.mat, 0.95)
## Sparsity = 0.95 (7 terms remaining) ~ 0.97 (17 terms remaining) seems acceptable
```

```
## Attach class label
alldata <- as.matrix(dt.mat.use)
alldata <- cbind(alldata, use$party)
colnames(alldata)[ncol(alldata)] <- "Class"
## Class=1 for Democrats, Class=2 for Republican
alldata <- as.data.frame(alldata)
```

```
alldata$Class <- as.factor(alldata$Class)
levels(alldata$Class) <- c("Democrats", "Republican")
```

Use 10-fold CV with 70% data for training, 30% data for testing

```
## Train-Test Split
set.seed(9000)
## 70% for training, 30% for testing
TrainingDataIndex <- createDataPartition(alldata$Class, p=0.7, list = FALSE)
trainingData <- alldata[TrainingDataIndex,]
testData <- alldata[-TrainingDataIndex,]
## 10-fold CV (cannot do repeatedcv due to CPU performance)
TrainingParameters <- trainControl(method = "cv", number = 10, classProbs = TRUE, summaryFunction = two
```

Classification using kNN (too many ties)

```
# ## Training
# fit <- train(Class ~ ., data = trainingData,
#             method = "knn",
#             trControl= TrainingParameters,
#             tuneGrid = expand.grid(k = seq(1, 10, length = 10)),
#             preProcess = c("scale","center"),
#             na.action = na.omit
# )
# fit
# fit$bestTune
#
# ## Testing
# pred <- predict(fit, testData)
#
# ## Evaluation
# confusionMatrix(pred, testData$Class)
#
# ## Rank terms by importance
# importance <- varImp(fit, scale=FALSE)
# plot(importance)
```

Classification using SVM

```
set.seed(9001)
## Using linear kernel

## Training
fit.LSVM <- train(Class ~ ., data = trainingData,
                  method = "svmLinear",
                  metric = "ROC",
                  trControl= TrainingParameters,
                  tuneGrid = expand.grid(C = seq(0.1, 1, length = 9)),
                  preProcess = c("scale","center"),
                  na.action = na.omit
)
```

```
## maximum number of iterations reached 0.001908255 0.001890649maximum number of iterations reached 0.0
```

```
fit.lSVM
```


```
## Support Vector Machines with Linear Kernel
##
## 2010 samples
## 10 predictor
## 2 classes: 'Democrats', 'Republican'
##
## Pre-processing: scaled (10), centered (10)
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 1809, 1810, 1809, 1808, 1809, 1808, ...
## Resampling results across tuning parameters:
##
## C ROC Sens Spec
## 0.1000 0.5216105 0.9980392 0.00000000
## 0.2125 0.5158934 0.9876114 0.01679965
## 0.3250 0.5496734 0.9876114 0.01679965
## 0.4375 0.5233665 0.9980392 0.00000000
## 0.5500 0.4952598 0.9921569 0.01041667
## 0.6625 0.5177286 0.9980392 0.00000000
## 0.7750 0.5553563 0.9980392 0.00000000
## 0.8875 0.5149111 0.9980392 0.00000000
## 1.0000 0.5243349 0.9980392 0.00000000
##
## ROC was used to select the optimal model using the largest value.
## The final value used for the model was C = 0.775.
```

```
fit.lSVM$bestTune
```

```
## C
## 7 0.775
```

```
## Testing
pred.lSVM <- predict(fit.lSVM, testData, type="prob")

## Evaluation
#confusionMatrix(pred, testData$Class)
result.roc <- roc(testData$Class, pred.lSVM$Democrats)
plot(result.roc, print.thres="best", print.thres.best.method="closest.topleft")
```



```
result.coords <- coords(result.roc, "best", best.method="closest.topleft", ret=c("threshold", "accuracy"))
print(result.coords)
```

```
##           threshold accuracy
## threshold 0.770811 0.3867596
```

```
# set.seed(9002)
# ## Using polynomial kernel
#
# ## Training
# ## Due to CPU performance, cannot apply a grid to tune parameters
# fit.pSVM <- train(Class ~ ., data = trainingData,
#                   method = "svmPoly",
#                   metric = "ROC",
#                   trControl= TrainingParameters,
#                   preProcess = c("scale","center"),
#                   na.action = na.omit
# )
# fit.pSVM
# fit.pSVM$bestTune
#
# ## Testing
# pred.pSVM <- predict(fit.pSVM, testData)
#
# ## Evaluation
```

```

# #confusionMatrix(pred, testData$Class)
# result.roc <- roc(testData$Class, pred.pSVM$Democrats)
# plot(result.roc, print.thres="best", print.thres.best.method="closest.topleft")
# result.coords <- coords(result.roc, "best", best.method="closest.topleft", ret=c("threshold", "accuracy"))
# print(result.coords)

```

```

set.seed(9003)
## Using radial basis kernel

## Training
## Due to CPU performance, cannot apply a grid to tune parameters
fit.rSVM <- train(Class ~ ., data = trainingData,
  method = "svmRadial",
  metric = "ROC",
  trControl= TrainingParameters,
  preProcess = c("scale","center"),
  na.action = na.omit
)
fit.rSVM

```

```

## Support Vector Machines with Radial Basis Function Kernel
##
## 2010 samples
## 10 predictor
## 2 classes: 'Democrats', 'Republican'
##
## Pre-processing: scaled (10), centered (10)
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 1810, 1808, 1809, 1810, 1809, 1809, ...
## Resampling results across tuning parameters:
##
## C ROC Sens Spec
## 0.25 0.5894976 0.9915372 0.07974291
## 0.50 0.5697093 0.9869833 0.09649823
## 1.00 0.5666541 0.9804728 0.12180851
##
## Tuning parameter 'sigma' was held constant at a value of 0.104267
## ROC was used to select the optimal model using the largest value.
## The final values used for the model were sigma = 0.104267 and C = 0.25.

```

```
fit.rSVM$bestTune
```

```

## sigma C
## 1 0.104267 0.25


```

```

## Testing
pred.rSVM <- predict(fit.rSVM, testData, type="prob")

## Evaluation
#confusionMatrix(pred, testData$Class)
result.roc <- roc(testData$Class, pred.rSVM$Democrats)
plot(result.roc, print.thres="best", print.thres.best.method="closest.topleft")

```



```
result.coords <- coords(result.roc, "best", best.method="closest.topleft", ret=c("threshold", "accuracy"))
print(result.coords)
```

```
##           threshold accuracy
## threshold 0.7820195 0.4982578
```

Classification using Naive Bayes


```
set.seed(9004)
## Training
fit.NB <- train(trainingData[, -ncol(trainingData)], trainingData$Class,
  method = "nb",
  metric = "ROC",
  preProcess=c("scale", "center"),
  trControl= TrainingParameters,
  na.action = na.omit
)
fit.NB
```

```
## Naive Bayes
##
## 2010 samples
## 10 predictor
## 2 classes: 'Democrats', 'Republican'
##
```

```
## Pre-processing: scaled (10), centered (10)
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 1809, 1809, 1808, 1810, 1809, 1808, ...
## Resampling results across tuning parameters:
##
##   usekernel  ROC      Sens      Spec
##   FALSE      0.6284958 0.8663908 0.2732713
##   TRUE       0.6362737 1.0000000 0.0000000
##
## Tuning parameter 'fL' was held constant at a value of 0
## Tuning
##   parameter 'adjust' was held constant at a value of 1
## ROC was used to select the optimal model using the largest value.
## The final values used for the model were fL = 0, usekernel = TRUE and adjust
##   = 1.
```

```
## Testing
pred.NB <- predict(fit.NB, testData, type="prob")

## Evaluation
#confusionMatrix(pred, testData$Class)
result.roc <- roc(testData$Class, pred.NB$Democrats)
plot(result.roc, print.thres="best", print.thres.best.method="closest.topleft")
```



```
result.coords <- coords(result.roc, "best", best.method="closest.topleft", ret=c("threshold", "accuracy"))
print(result.coords)
```

```
##           threshold accuracy
## threshold 0.9827756 0.3507549
```


Classification using Logistic Regression

```
set.seed(9005)
## Training
fit.LR <- train(Class ~ ., data = trainingData,
                method = "glm",
                metric = "ROC",
                preProcess=c("scale", "center"),
                trControl= TrainingParameters,
                na.action = na.omit
)
fit.LR
```

```
## Generalized Linear Model
##
## 2010 samples
## 10 predictor
## 2 classes: 'Democrats', 'Republican'
##
## Pre-processing: scaled (10), centered (10)
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 1809, 1810, 1808, 1809, 1809, 1810, ...
## Resampling results:
##
##      ROC      Sens      Spec
## 0.6360117 0.9823954 0.05066489
```

```
## Testing
pred.LR <- predict(fit.LR, testData, type="prob")

## Evaluation
#confusionMatrix(pred, testData$Class)
result.roc <- roc(testData$Class, pred.LR$Democrats)
plot(result.roc, print.thres="best", print.thres.best.method="closest.topleft")
```

```
result.coords <- coords(result.roc, "best", best.method="closest.topleft", ret=c("threshold", "accuracy"))
print(result.coords)
```

```
##           threshold accuracy
## threshold 0.7647984 0.3577236
```

Classification using Decision Tree


```
set.seed(9006)
## Training
fit.DT <- train(Class ~ ., data = trainingData,
  method = "rpart",
  metric = "ROC",
  preProcess=c("scale","center"),
  trControl= TrainingParameters,
  na.action = na.omit
)
fit.DT
```

```
## CART
##
## 2010 samples
## 10 predictor
## 2 classes: 'Democrats', 'Republican'
##
```

```
## Pre-processing: scaled (10), centered (10)
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 1808, 1810, 1808, 1809, 1809, 1809, ...
## Resampling results across tuning parameters:
##
##      cp          ROC      Sens      Spec
## 0.01260504 0.5258730 0.9771878 0.08204787
## 0.01470588 0.5167761 0.9791529 0.06121454
## 0.01785714 0.5090492 0.9817673 0.04228723
##
## ROC was used to select the optimal model using the largest value.
## The final value used for the model was cp = 0.01260504.
```

```
## Testing
pred.DT <- predict(fit.DT, testData, type="prob")

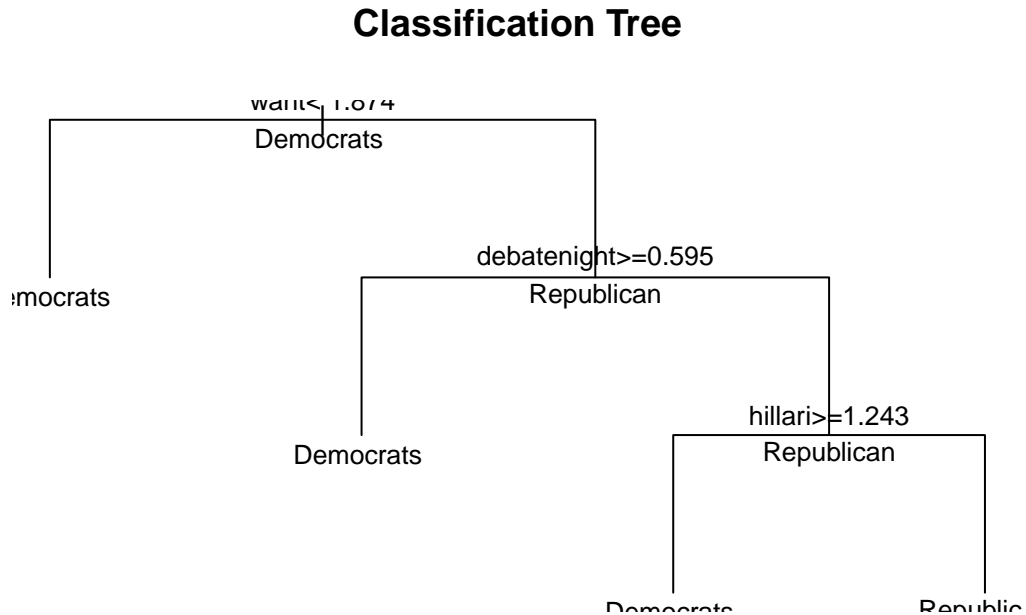
## Evaluation
#confusionMatrix(pred, testData$Class)
result.roc <- roc(testData$Class, pred.DT$Democrats)
plot(result.roc, print.thres="best", print.thres.best.method="closest.topleft")
```



```
result.coords <- coords(result.roc, "best", best.method="closest.topleft", ret=c("threshold", "accuracy"))
print(result.coords)
```

```
##           threshold accuracy
## threshold 0.7564035 0.2299652
```

```
plot(fit.DT$finalModel, uniform=TRUE, main="Classification Tree")
text(fit.DT$finalModel, use.n.=TRUE, all=TRUE, cex=.8)
```



Classification using AdaBoost


```
set.seed(9007)
## Training
fit.ADA <- train(trainingData[, -ncol(trainingData)], trainingData$Class,
  method = "ada",
  metric = "ROC",
  preProcess=c("scale", "center"),
  trControl= TrainingParameters,
  na.action = na.omit
)
fit.ADA
```

```
## Boosted Classification Trees
##
## 2010 samples
## 10 predictor
## 2 classes: 'Democrats', 'Republican'
##
## Pre-processing: scaled (10), centered (10)
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 1809, 1808, 1808, 1808, 1809, 1809, ...
```

```
## Resampling results across tuning parameters:
##
##   maxdepth  iter  ROC      Sens      Spec
##   1         50   0.6121298 1.0000000 0.000000000
##   1         100  0.6306340 0.9980435 0.004166667
##   1         150  0.6368190 0.9980435 0.008333333
##   2         50   0.6385099 0.9856634 0.050310284
##   2         100  0.6378876 0.9811094 0.084086879
##   2         150  0.6383334 0.9785035 0.098714539
##   3         50   0.6412266 0.9837026 0.102703901
##   3         100  0.6414884 0.9830490 0.109042553
##   3         150  0.6430482 0.9817503 0.111170213
##
## Tuning parameter 'nu' was held constant at a value of 0.1
## ROC was used to select the optimal model using the largest value.
## The final values used for the model were iter = 150, maxdepth = 3 and nu = 0.1.
```

```
## Testing
pred.ADA <- predict(fit.ADA, testData, type="prob")

## Evaluation
#confusionMatrix(pred, testData$Class)
result.roc <- roc(testData$Class, pred.ADA$Democrats)
plot(result.roc, print.thres="best", print.thres.best.method="closest.topleft")
```



```
result.coords <- coords(result.roc, "best", best.method="closest.topleft", ret=c("threshold", "accuracy")  
print(result.coords)
```

```
##           threshold  accuracy  
## threshold 0.7564355 0.3507549
```