

# classification\_Xinming

Xinming Liu

11/27/2020

Set A - Debate 1

```
## SetA Preprocessing
users <- read.csv("user_setA/users.csv", stringsAsFactors = FALSE)
#head(a.users)
tweets1 <- read.csv("user_setA/tweets_debate1.csv", stringsAsFactors = FALSE)
tweets2 <- read.csv("user_setA/tweets_debate2.csv", stringsAsFactors = FALSE)
tweets3 <- read.csv("user_setA/tweets_debate3.csv", stringsAsFactors = FALSE)
tweets4 <- read.csv("user_setA/tweets_debateVP.csv", stringsAsFactors = FALSE)
tweets <- rbind(tweets1, tweets2, tweets3, tweets4)
#head(tweets)
tweets<-tweets%>%
  mutate(userID=as.numeric((userID)))
data <- data.table(users, key="userID")[
  data.table(tweets, key="userID"),
  allow.cartesian=TRUE
]
data <- subset(data, party=="D" | party=="R")
data <- subset(data, state_code=="FL")
#head(data)
use <- data.frame(data$text, factor(data$party), stringsAsFactors = FALSE)
colnames(use) <- c("text", "party")
#head(use)
```

```
## Create DocumentTermMatrix
corpus <- Corpus(VectorSource(use$text))
corpus = clean_corpus(corpus)
#td.mat = TermDocumentMatrix(corpus)
dt.mat = DocumentTermMatrix(corpus)
## dt.mat is not a matrix here
```

```
## Feature words extraction (this may create NA values) due to limited memory
dt.mat.use = removeSparseTerms(dt.mat, 0.95)
## Sparsity = 0.95 (7 terms remaining) ~ 0.97 (17 terms remaining) seems acceptable
```

```
## Attach class label
alldata <- as.matrix(dt.mat.use)
alldata <- cbind(alldata, use$party)
colnames(alldata)[ncol(alldata)] <- "Class"
## Class=1 for Democrats, Class=2 for Republican
alldata <- as.data.frame(alldata)
```

```
alldata$Class <- as.factor(alldata$Class)
levels(alldata$Class) <- c("Democrats", "Republican")
```

Use 10-fold CV with 70% data for training, 30% data for testing

```
## Train-Test Split
set.seed(9000)
## 70% for training, 30% for testing
TrainingDataIndex <- createDataPartition(alldata$Class, p=0.7, list = FALSE)
trainingData <- alldata[TrainingDataIndex,]
testData <- alldata[-TrainingDataIndex,]
## 10-fold CV (cannot do repeatedcv due to CPU performance)
TrainingParameters <- trainControl(method = "cv", number = 10, classProbs = TRUE, summaryFunction = two
```

Classification using kNN (too many ties)

```
# ## Training
# fit <- train(Class ~ ., data = trainingData,
#             method = "knn",
#             trControl= TrainingParameters,
#             tuneGrid = expand.grid(k = seq(1, 10, length = 10)),
#             preProcess = c("scale","center"),
#             na.action = na.omit
# )
# fit
# fit$bestTune
#
# ## Testing
# pred <- predict(fit, testData)
#
# ## Evaluation
# confusionMatrix(pred, testData$Class)
#
# ## Rank terms by importance
# importance <- varImp(fit, scale=FALSE)
# plot(importance)
```

Classification using SVM

```
set.seed(9001)
## Using linear kernel

## Training
fit.LSVM <- train(Class ~ ., data = trainingData,
                  method = "svmLinear",
                  metric = "ROC",
                  trControl= TrainingParameters,
                  tuneGrid = expand.grid(C = seq(0.1, 1, length = 9)),
                  preProcess = c("scale","center"),
                  na.action = na.omit
)
```

```
## maximum number of iterations reached 0.0009052221 0.0008988652maximum number of iterations reached 0
```

```
fit.lSVM
```


```
## Support Vector Machines with Linear Kernel
##
## 563 samples
## 6 predictor
## 2 classes: 'Democrats', 'Republican'
##
## Pre-processing: scaled (6), centered (6)
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 507, 507, 507, 507, 507, 506, ...
## Resampling results across tuning parameters:
##
##  C          ROC          Sens  Spec
##  0.1000  0.4957998  1      0
##  0.2125  0.4745358  1      0
##  0.3250  0.4898983  1      0
##  0.4375  0.5328424  1      0
##  0.5500  0.4858082  1      0
##  0.6625  0.4530615  1      0
##  0.7750  0.4991108  1      0
##  0.8875  0.5218775  1      0
##  1.0000  0.5133029  1      0
##
## ROC was used to select the optimal model using the largest value.
## The final value used for the model was C = 0.4375.
```

```
fit.lSVM$bestTune
```

```
##          C
## 4 0.4375
```

```
## Testing
pred.lSVM <- predict(fit.lSVM, testData, type="prob")

## Evaluation
#confusionMatrix(pred, testData$Class)
result.roc <- roc(testData$Class, pred.lSVM$Democrats)
plot(result.roc, print.thres="best", print.thres.best.method="closest.topleft")
```



```
result.coords <- coords(result.roc, "best", best.method="closest.topleft", ret=c("threshold", "accuracy"))
print(result.coords)
```

```
##           threshold accuracy
## threshold 0.7122004 0.5020921
```

```
# set.seed(9002)
# ## Using polynomial kernel
#
# ## Training
# ## Due to CPU performance, cannot apply a grid to tune parameters
# fit.pSVM <- train(Class ~ ., data = trainingData,
#                   method = "svmPoly",
#                   metric = "ROC",
#                   trControl= TrainingParameters,
#                   preProcess = c("scale","center"),
#                   na.action = na.omit
# )
# fit.pSVM
# fit.pSVM$bestTune
#
# ## Testing
# pred.pSVM <- predict(fit.pSVM, testData)
#
# ## Evaluation
```

```

# #confusionMatrix(pred, testData$Class)
# result.roc <- roc(testData$Class, pred.pSVM$Democrats)
# plot(result.roc, print.thres="best", print.thres.best.method="closest.topleft")
# result.coords <- coords(result.roc, "best", best.method="closest.topleft", ret=c("threshold", "accuracy"))
# print(result.coords)

```

```

set.seed(9003)
## Using radial basis kernel

## Training
## Due to CPU performance, cannot apply a grid to tune parameters
fit.rSVM <- train(Class ~ ., data = trainingData,
  method = "svmRadial",
  metric = "ROC",
  trControl= TrainingParameters,
  preProcess = c("scale","center"),
  na.action = na.omit
)
fit.rSVM

```

```

## Support Vector Machines with Radial Basis Function Kernel
##
## 563 samples
## 6 predictor
## 2 classes: 'Democrats', 'Republican'
##
## Pre-processing: scaled (6), centered (6)
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 507, 506, 507, 508, 507, 507, ...
## Resampling results across tuning parameters:
##
##  C      ROC      Sens      Spec
##  0.25  0.4787260  0.9923077  0.005882353
##  0.50  0.5353951  0.9820513  0.005882353
##  1.00  0.5433448  0.9897436  0.005555556
##
## Tuning parameter 'sigma' was held constant at a value of 0.1043378
## ROC was used to select the optimal model using the largest value.
## The final values used for the model were sigma = 0.1043378 and C = 1.

```

```
fit.rSVM$bestTune
```

```

##      sigma C
## 3 0.1043378 1


```

```

## Testing
pred.rSVM <- predict(fit.rSVM, testData, type="prob")

## Evaluation
#confusionMatrix(pred, testData$Class)
result.roc <- roc(testData$Class, pred.rSVM$Democrats)
plot(result.roc, print.thres="best", print.thres.best.method="closest.topleft")

```



```
result.coords <- coords(result.roc, "best", best.method="closest.topleft", ret=c("threshold", "accuracy"))
print(result.coords)
```

```
##           threshold accuracy
## threshold 0.6872721 0.4853556
```

Classification using Naive Bayes


```
set.seed(9004)
## Training
fit.NB <- train(trainingData[, -ncol(trainingData)], trainingData$Class,
  method = "nb",
  metric = "ROC",
  preProcess=c("scale", "center"),
  trControl= TrainingParameters,
  na.action = na.omit
)
fit.NB
```

```
## Naive Bayes
##
## 563 samples
## 6 predictor
## 2 classes: 'Democrats', 'Republican'
##
```

```
## Pre-processing: scaled (6), centered (6)
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 507, 507, 506, 507, 508, 506, ...
## Resampling results across tuning parameters:
##
##   usekernel  ROC      Sens      Spec
##   FALSE      0.6128860 0.3142375 0.8866013
##   TRUE       0.6080343 0.4427800 0.7294118
##
## Tuning parameter 'fL' was held constant at a value of 0
## Tuning
##   parameter 'adjust' was held constant at a value of 1
## ROC was used to select the optimal model using the largest value.
## The final values used for the model were fL = 0, usekernel = FALSE and adjust
##   = 1.
```

```
## Testing
pred.NB <- predict(fit.NB, testData, type="prob")

## Evaluation
#confusionMatrix(pred, testData$Class)
result.roc <- roc(testData$Class, pred.NB$Democrats)
plot(result.roc, print.thres="best", print.thres.best.method="closest.topleft")
```



```
result.coords <- coords(result.roc, "best", best.method="closest.topleft", ret=c("threshold", "accuracy"))
print(result.coords)
```

```
##           threshold accuracy
## threshold 0.510868 0.5857741
```

Classification using Logistic Regression


```
set.seed(9005)
## Training
fit.LR <- train(Class ~ ., data = trainingData,
               method = "glm",
               metric = "ROC",
               preProcess=c("scale", "center"),
               trControl= TrainingParameters,
               na.action = na.omit
)
fit.LR
```

```
## Generalized Linear Model
##
## 563 samples
## 6 predictor
## 2 classes: 'Democrats', 'Republican'
##
## Pre-processing: scaled (6), centered (6)
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 508, 506, 506, 506, 507, 507, ...
## Resampling results:
##
##      ROC      Sens      Spec
## 0.606487 0.9717949 0.05196078
```

```
## Testing
pred.LR <- predict(fit.LR, testData, type="prob")

## Evaluation
#confusionMatrix(pred, testData$Class)
result.roc <- roc(testData$Class, pred.LR$Democrats)
plot(result.roc, print.thres="best", print.thres.best.method="closest.topleft")
```





```
result.coords <- coords(result.roc, "best", best.method="closest.topleft", ret=c("threshold", "accuracy"))
print(result.coords)
```

```
##           threshold accuracy
## threshold 0.7477121 0.5857741
```

Classification using Decision Tree


```
set.seed(9006)
## Training
fit.DT <- train(Class ~ ., data = trainingData,
  method = "rpart",
  metric = "ROC",
  preProcess=c("scale","center"),
  trControl= TrainingParameters,
  na.action = na.omit
)
fit.DT
```

```
## CART
##
## 563 samples
## 6 predictor
## 2 classes: 'Democrats', 'Republican'
##
```

```
## Pre-processing: scaled (6), centered (6)
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 506, 507, 506, 507, 507, 506, ...
## Resampling results across tuning parameters:
##
##      cp          ROC      Sens      Spec
## 0.000000000 0.5459914 0.9973684 0.011437908
## 0.001142857 0.5459914 0.9973684 0.011437908
## 0.002285714 0.5227008 0.9973684 0.005555556
##
## ROC was used to select the optimal model using the largest value.
## The final value used for the model was cp = 0.001142857.
```

```
## Testing
pred.DT <- predict(fit.DT, testData, type="prob")

## Evaluation
#confusionMatrix(pred, testData$Class)
result.roc <- roc(testData$Class, pred.DT$Democrats)
plot(result.roc, print.thres="best", print.thres.best.method="closest.topleft")
```

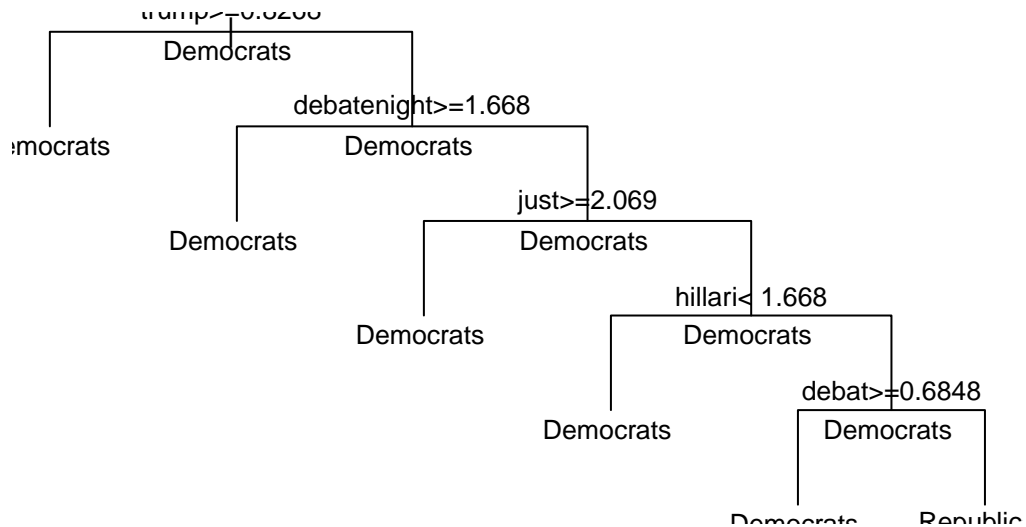


```
result.coords <- coords(result.roc, "best", best.method="closest.topleft", ret=c("threshold", "accuracy"))
print(result.coords)
```

```
##          threshold accuracy
## threshold 0.6754386 0.5857741
```

```
plot(fit.DT$finalModel, uniform=TRUE, main="Classification Tree")
text(fit.DT$finalModel, use.n.=TRUE, all=TRUE, cex=.8)
```

## Classification Tree



Classification using AdaBoost


```
set.seed(9007)
## Training
fit.ADA <- train(trainingData[, -ncol(trainingData)], trainingData$Class,
  method = "ada",
  metric = "ROC",
  preProcess=c("scale", "center"),
  trControl= TrainingParameters,
  na.action = na.omit
)
fit.ADA
```

```
## Boosted Classification Trees
##
## 563 samples
## 6 predictor
## 2 classes: 'Democrats', 'Republican'
##
## Pre-processing: scaled (6), centered (6)
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 507, 507, 506, 506, 507, 507, ...
```

```
## Resampling results across tuning parameters:
##
##   maxdepth  iter  ROC      Sens      Spec
##   1         50   0.5784609 1.0000000 0.000000000
##   1         100  0.5987609 1.0000000 0.000000000
##   1         150  0.6035772 1.0000000 0.000000000
##   2         50   0.6015338 1.0000000 0.000000000
##   2         100  0.6083345 0.9948718 0.005882353
##   2         150  0.6089374 0.9844804 0.022549020
##   3         50   0.6120182 0.9974359 0.000000000
##   3         100  0.6137247 0.9818489 0.028104575
##   3         150  0.6141977 0.9818489 0.028104575
##
## Tuning parameter 'nu' was held constant at a value of 0.1
## ROC was used to select the optimal model using the largest value.
## The final values used for the model were iter = 150, maxdepth = 3 and nu = 0.1.
```

```
## Testing
pred.ADA <- predict(fit.ADA, testData, type="prob")

## Evaluation
#confusionMatrix(pred, testData$Class)
result.roc <- roc(testData$Class, pred.ADA$Democrats)
plot(result.roc, print.thres="best", print.thres.best.method="closest.topleft")
```



```
result.coords <- coords(result.roc, "best", best.method="closest.topleft", ret=c("threshold", "accuracy")  
print(result.coords)
```

```
##           threshold accuracy  
## threshold 0.7963398 0.5983264
```